Project Report

Decision Support System for Smart Irrigation

Qatra

Group Members :

Azza Jouini, ID: 241FTB0765 Azza.JOUINI@esprit.tn

Hanna Hmouda, ID: 241FTB0165 Hanna.HMOUDA@esprit.tn

Salsabil Ben ElHadj, ID: 241FTB1184 BenElhadj.SALSABIL@esprit.tn

Firas Ben Salem, ID: 241MTB1570 MohamedFiras.BENSALEM@esprit.tn

Salah Chebbi, ID: 241MTB1851 Salah.CHEBBI@esprit.tn

Raed Meddeb, ID: 231MTB1459 Raed.MEDDEB@esprit.tn

Table of Contents

Project Overview

Agriculture is one of the most water-intensive sectors, especially in regions facing water scarcity such as Tunisia. Efficient water management has become a critical challenge for farmers who must balance productivity with sustainability. Traditional irrigation methods often rely on fixed schedules or manual decisions, leading to water waste and reduced efficiency.

This project aims to develop a Decision Support System (DSS) named Qatra that assists farmers in optimizing water usage through data analysis and machine learning techniques.

Phase 1: Problem Definition & Requirements

1. Study of the Existing Solutions: Analysis and Comparison

The project began with a discovery phase aimed at identifying the gap in the current agricultural market.

System	Focus Area	Key Features	Limitations / The Gap
FAO CropWat	Irrigation Planning	Calculates water requirements based on climate, soil, and crop data.	Requires manual data entry; not real-time or automated for daily farm management.
Agricolus	Precision Farming	Uses satellite imagery and weather forecasts to map field needs.	Expensive subscription models; can be too complex for smallholder farmers.
APSIM / DSSAT	Crop Simulation	Predicts long-term yield based on various management scenarios.	Aimed at researchers/scientists rather than daily decision-making for farmers.

2. Problem Statement

- **Current State**: Most local farmers rely on manual or timer-based irrigation, which does not account for actual soil health or weather changes.

- **The Problematic**:
    - **Resource Waste**: Over-irrigation leads to soil degradation and depletion of water reserves.
    - **Economic Impact**: Inefficient water use increases costs and can reduce crop yields if plants are under-irrigated.

3. Proposed Solution

The proposed solution, **Qatra**, is an integrated Decision Support System designed to transition agricultural water management from traditional fixed schedules to data-driven precision irrigation. **Qatra** aims to optimize water usage thus improving agricultural efficiency and addressing water scarcity issues.

The system architecture is divided into two primary modules:

3.1 Analytical Module (Data Warehousing & BI)

This component focuses on historical data exploration to identify trends and inefficiencies.

- **Process**: Using Talend as an ETL tool, raw agricultural data (Soil Type, Crop Type, Water Usage) is extracted, transformed, and loaded into a centralized Data Warehouse designed with a Star Schema.
- **Output**: Interactive Power BI dashboards provide stakeholders with visual insights into water consumption patterns and yield performance across different seasons and irrigation methods.

3.2 Predictive Module (Machine Learning)

This component serves as the core "intelligence" of the system, shifting the focus from historical analysis to actionable prediction.

- **Process**: Using Python and the Scikit-learn library, the system first applies unsupervised learning to engineer the target variable:

    - Dimensionality Reduction: Principal Component Analysis (PCA) is applied to the preprocessed features to reduce noise, remove multicollinearity, and prepare the data for clustering.
    - Clustering & Label Engineering: Multiple unsupervised algorithms are tested — including K-Means, DBSCAN, and Agglomerative Hierarchical Clustering. Each is evaluated using the Silhouette Score, Davies-Bouldin Index, and the Elbow Method. The best-performing algorithm is selected to assign irrigation efficiency cluster labels (e.g., Efficient vs. Inefficient) to each farm record.
    - Supervised Classification: The generated cluster labels are then used as the target variable to train and compare supervised classifiers (Random Forest, SVM, KNN), enabling interpretable and actionable irrigation recommendations.
- **Output**: The module generates recommendations for irrigation strategy based on current farm variables, ensuring that water is applied only when and where it is needed to maximize crop yield while minimizing waste.

4. Competitive Differentiation

While professional agricultural tools exist, Qatra distinguishes itself by bridging the gap between historical oversight and real-time precision through the following core advantages:

- **Cost & Accessibility**: Many existing DSS require high-subscription fees or expensive hardware. Qatra focuses on being affordable for small-medium farmers.
- **Integration**: Unlike standalone apps, Qatra combines historical analysis (Data Warehouse/Power BI) with predictive intelligence (Python ML), offering a complete end-to-end decision support pipeline.

5. Requirement Analysis

5.1 Functional Requirements

- The system must collect and store data on soil type, crop type, irrigation method, and water usage.

- The ML pipeline must first apply unsupervised clustering to engineer the target variable, then train a supervised classifier to generate actionable irrigation recommendations.
- The dashboard must be interactive for the farmer to interpret yield trends.

5.2 Non- Functional Requirements

- **Data Privacy:** Ensuring farm-specific data (Farm_ID) is handled securely.
- **Accessibility:** Designing the Power BI dashboard to be user-friendly and simple enough for users who may not be data experts.
- **Reliability**: Provide accurate and reliable predictions.

5.3 Technical Requirements

- **Data Source:** The Agriculture CSV dataset (5,000 records, 10 features).
- **ETL Layer:** Talend processing the data.
- **Storage Layer**: A Relational DBMS for the Data Warehouse (Star Schema)
- **Presentation Layer:** Power BI Dashboards for visualization.
- **Modeling**: Python (Jupyter Lab) implementing a two-stage ML pipeline: unsupervised clustering (K-Means, DBSCAN, Agglomerative) with PCA preprocessing for label engineering, followed by supervised classification (Random Forest, SVM, KNN) for prediction.

Insight: With 5,000 records and 10 features, the dataset is well-suited for robust unsupervised and supervised learning. PCA is applied prior to clustering to reduce dimensionality and ensure the algorithms converge meaningfully. The best clustering algorithm is selected based on quantitative evaluation metrics before its labels are used to train the downstream classifier.

6. Dataset Description

This section provides a clear overview of the data we will be using to build the Decision Support System.

- **Dataset Name**: Agriculture Dataset.
- **Nature of the Data**: The dataset contains 5000 records representing various farm configurations, environmental conditions, and resource usage metrics. There are no missing values across all 10 columns.
- **Column Breakdown:**
    - **Farm_ID**: A unique identifier for each agricultural plot (e.g., SYN_0001).
    - **Crop_Type**: The specific crop planted: Potato, Rice, Carrot, Tomato, Soybean, Barley, Cotton, Wheat, Maize, or Sugarcane (10 crop types).
    - **Farm_Area(acres)**: The size of the farm in acres (range: 10 to 500).
    - **Irrigation_Type**: The method used for watering, including Sprinkler, Manual, Flood, Rain-fed, and Drip.
    - **Fertilizer_Used(tons)**: The quantity of fertilizer applied.
    - **Pesticide_Used(kg)**: The quantity of pesticide applied to the crops.
    - **Yield(tons)**: The total agricultural output produced.
    - **Soil_Type**: The classification of the soil, such as Loamy, Peaty, Silty, Clay, or Sandy.
    - **Season**: The specific growing season (Kharif, Rabi, or Zaid).
    - **Water_Usage(cubic meters)**: The volume of water consumed, which is a critical feature for your optimization goals. Values range from approximately 5,000 to 95,000 cubic meters.

7. Project Success Metrics: Key Performance Indicators (KPIs)

To evaluate the efficiency and impact of the **Qatra** Decision Support System, the following Key Performance Indicators (KPIs) have been established:

- **Water Savings Percentage**: This metric measures the reduction in the volume of water consumed (in cubic meters) by comparing the data-driven irrigation recommendations against traditional fixed-schedule baselines. The goal is to minimize waste without compromising soil health.
- **Yield Stability**: A critical indicator to ensure that optimized water usage does not lead to under-irrigation. Success is defined by maintaining or increasing the total agricultural output (in tons) while using fewer resources.
- **Model Predictive Accuracy**: This measures the technical reliability of the Python-based Machine Learning models. Success is determined by the model's precision and its ability to correctly categorize irrigation needs, ensuring that the DSS provides trustworthy advice to the farmer.

8. Project Implementation Roadmap

- **Phase 1:** Specifications – Finalizing the project scope and technical constraints.
- **Phase 2:** Pre-processing – Data cleaning, encoding, scaling, normalization, PCA dimensionality reduction, and unsupervised clustering for target variable engineering.
- **Phase 3:** Warehouse Design – Designing Fact and Dimension tables.
- **Phase 4:** Data Feeding – Executing ETL jobs with Talend.
- **Phase 5:** Visualization – Data exploration and reporting via Power BI.
- **Phase 6:** Data Modeling – Supervised classification and evaluation using Python-based ML, building on the cluster labels generated in Phase 2.

Phase 2: DATA PRE-PROCESSING

1. Environment Setup & Data Acquisition

The first step of the pre-processing phase involved setting up the computational environment and loading the raw data. We utilized the **Pandas** and **NumPy** libraries for data manipulation, and **Matplotlib/Seaborn** for future visualizations.

The "Agriculture Dataset" was loaded from a CSV file into a Pandas DataFrame named dataset. This structure allows for efficient indexing and analysis of the 5000 farm records that form the foundation of the Qatra Decision Support System.

```
# 1. Import the Libraries and load the dataset
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
# Load the dataset
dataset = pd.read_csv('smart_agriculture_features.csv')
```

2. Exploratory Data Analysis (EDA): Structural Overview

Before applying any transformations, it was necessary to understand the structure, quality, and statistical properties of the dataset. We executed several descriptive functions to audit the data.

- **Data Preview (head)**: We displayed the first 10 rows to display an overview of the dataset.

```
# 2.1 Data Preview
dataset.head(10)
```

	Farm_ID	Crop_Type	Farm_Area(acres)	Irrigation_Type	Fertilizer_Used(tons)	Pesticide_Used(kg)	Yield(tons)	Soil_Type	Season	Water_Usage(cubic meters)
0	SYN_0001	Potato	421.36	Drip	8.24	1.84	44.73	Peaty	Kharif	30106.65
1	SYN_0002	Rice	353.41	Rain-fed	6.30	2.49	34.39	Silty	Zaid	20471.93
2	SYN_0003	Potato	260.89	Flood	3.92	4.39	26.65	Silty	Kharif	28844.32
3	SYN_0004	Rice	368.83	Manual	1.29	3.98	34.87	Silty	Rabi	26743.29
4	SYN_0005	Potato	173.73	Sprinkler	2.36	2.87	38.25	Silty	Zaid	47760.90
5	SYN_0006	Carrot	83.44	Drip	6.99	1.70	48.90	Silty	Rabi	17255.42
6	SYN_0007	Carrot	296.71	Drip	9.31	0.79	32.67	Loamy	Rabi	10945.92
7	SYN_0008	Tomato	372.93	Manual	6.42	4.51	46.75	Sandy	Rabi	36667.05
8	SYN_0009	Soybean	162.13	Sprinkler	9.88	3.84	47.09	Sandy	Kharif	19862.61
9	SYN_0010	Barley	60.84	Sprinkler	4.47	3.36	38.50	Peaty	Kharif	41171.53

- **Dataset Dimensions (shape)**: The output confirmed a total of **5000 records** and **10 columns**, providing a clear boundary for our project scope.

```
# 2.2 Dataset Dimensions
dataset.shape
```

```
(5000, 10)
```

- **Integrity Check (isnull().sum)**: A critical audit for missing values was performed. The result showed **0 null values** across all columns, confirming that the dataset was complete and ready for processing without the need for imputation.

```
[5]: dataset.isnull().sum()
```

```
[5]: Farm_ID                  0
     Crop_Type                0
     Farm_Area(acres)         0
     Irrigation_Type          0
     Fertilizer_Used(tons)    0
     Pesticide_Used(kg)       0
     Yield(tons)              0
     Soil_Type                0
     Season                   0
     Water_Usage(cubic meters)  0
     dtype: int64
```

- **Structural Overview (info)**: This provided the data types for each column, confirming that five columns are numerical (float64) and five are categorical (object).

```
# 2.4 Structural Overview - Data Types
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 10 columns):
 #   Column                      Non-Null Count  Dtype
---  ------                      --------------  -----
 0   Farm_ID                     5000 non-null   object
 1   Crop_Type                   5000 non-null   object
 2   Farm_Area(acres)            5000 non-null   float64
 3   Irrigation_Type             5000 non-null   object
 4   Fertilizer_Used(tons)       5000 non-null   float64
 5   Pesticide_Used(kg)          5000 non-null   float64
 6   Yield(tons)                 5000 non-null   float64
 7   Soil_Type                   5000 non-null   object
 8   Season                      5000 non-null   object
 9   Water_Usage(cubic meters)   5000 non-null   float64
dtypes: float64(5), object(5)
memory usage: 390.8+ KB
```

- **Statistical Profiling (describe)**:
  - **Numerical Data**: Water_Usage ranges from ~5,000 to ~95,000 cubic meters with a mean of ~58,706, while Yield ranges from 1 to 55 tons. This confirmed a significant scale difference between features, justifying the need for feature scaling.

```
# 2.5 Statistical Profiling - Numerical Columns
dataset.describe()
```

	Farm_Area(acres)	Fertilizer_Used(tons)	Pesticide_Used(kg)	Yield(tons)	Water_Usage(cubic meters)
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	257.128142	5.305568	2.563260	22.338618	58705.761954
std	142.110064	2.740247	1.409917	15.891496	26286.333903
min	10.080000	0.500000	0.100000	1.000000	5005.300000
25%	133.325000	2.947500	1.357500	8.897500	36802.795000
50%	260.365000	5.400000	2.580000	16.880000	66892.375000
75%	382.932500	7.650000	3.790000	35.925000	81028.695000
max	499.990000	10.000000	5.000000	55.000000	94987.980000

  - **Categorical Data**: Using include=['object'], we identified the diversity of the data. We found **10 unique crop types**, **5 irrigation methods**, and **5 soil types**, with Barley (516), Drip irrigation (1,008), and Silty soil (1,056) being the most frequent categories.

```
# 2.6 Statistical Profiling - Categorical Columns
dataset.describe(include=['object'])
```

	Farm_ID	Crop_Type	Irrigation_Type	Soil_Type	Season
count	5000	5000	5000	5000	5000
unique	5000	10	5	5	3
top	SYN_5000	Barley	Drip	Silty	Kharif
freq	1	516	1008	1056	1677

- **Duplicates Check:** We executed a check for duplicate rows within the dataset. The audit confirmed 0 duplicate rows, ensuring each farm configuration is unique and contributes equally to the learning process.

```
# 2.7 Duplicates Check
print(f"Number of duplicate rows: {dataset.duplicated().sum()}")
```

```
Number of duplicate rows: 0
```

3. Visual Exploratory Data Analysis (EDA)

3.1 Distribution of Numerical Variables (Histograms)

To understand the underlying patterns and statistical health of the agricultural data, we generated histograms for all quantitative features. This step is essential to identify data distribution patterns and confirm the need for feature scaling.

- **Uniform Distribution:** Features like Farm_Area(acres), Fertilizer_Used(tons), and Pesticide_Used(kg) show a relatively uniform distribution across the 5000 records. This suggests that our dataset covers a wide variety of farm configurations ideal for training a robust model.

- **Skewed Distribution:** Yield(tons) shows a right-skewed distribution — most farms produce lower yields, with fewer high-yield farms. This is an important observation regarding real-world agricultural productivity.
- **Scale Variance:** A critical observation from these plots is the difference in magnitude. Yield values are clustered under 10 units, while Water_Usage extends beyond 80,000 units. This visual evidence confirms that we **must** apply Feature Scaling (Standardization) in later steps to prevent Water_Usage from overpowering other features during model training.

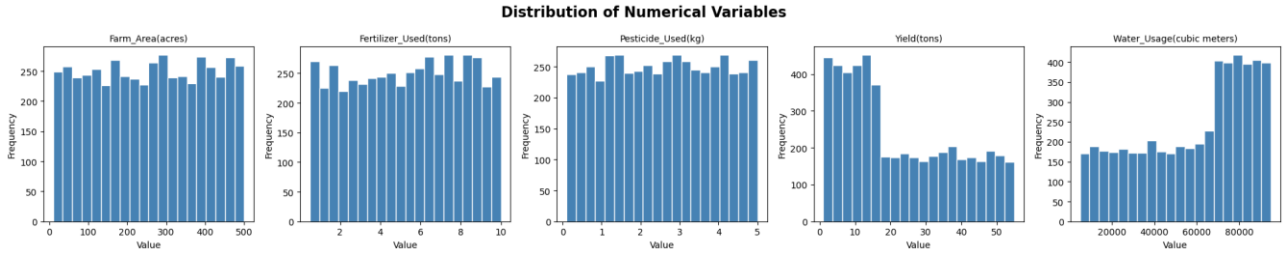

3.2 Outlier Detection

After analyzing the distributions, we performed a rigorous outlier detection using horizontal box plots. Due to the significant variance in scales between variables, we grouped the numerical columns into three separate plots to ensure each visualization remained interpretable and clear.

- **Group 1 - Farm_Area:** Wide and symmetric distribution spanning 10-500 acres with the median around 250. No outlier points visible beyond the whiskers.
- **Group 2 - Agricultural Inputs & Yield:** Yield(tons) spans a wide range with a symmetric box. Pesticide_Used and Fertilizer_Used are both compact and symmetric with narrow IQR ranges. Crucially, no outlier dots are visible beyond any whisker across all three features.
- **Group 3 - Water_Usage:** Wide range spanning approximately 5,000 to 95,000 cubic meters, symmetric around a median of approximately 65,000. Again, no outliers visible beyond the whiskers.

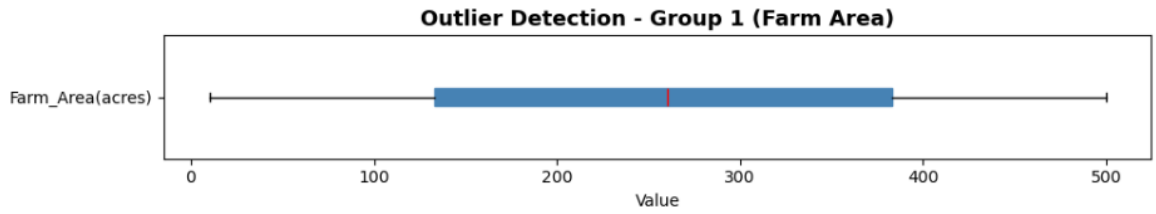

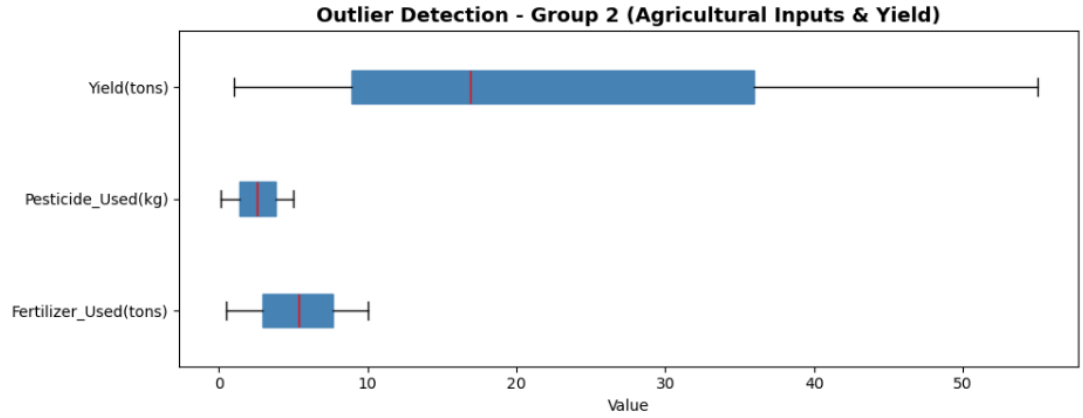

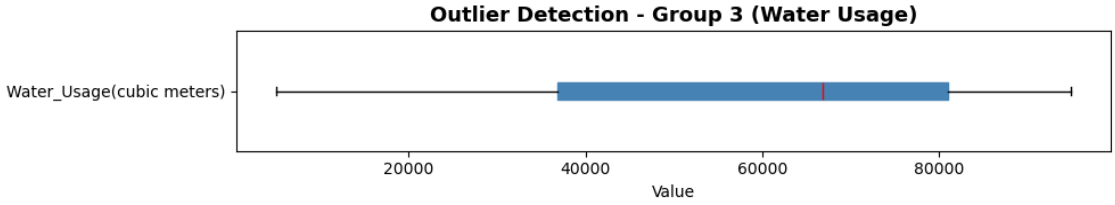

The dataset is remarkably clean with no statistical outliers across any of the five numerical features. This means we can proceed directly to modeling without any outlier removal or treatment.

3.3 Relationship Analysis (Correlation Heatmap)

To conclude the visual exploration, we generated a Pearson correlation heatmap for all numerical features. This matrix quantifies the linear relationship between variables, with values ranging from -1 (perfect negative correlation) to +1 (perfect positive correlation).

- **Dominant Finding - Yield vs Water_Usage (-0.72):** The most striking observation is the strong negative correlation between Yield and Water_Usage. This means farms that use more water tend to produce lower yields, and farms that use less water tend to produce higher yields. This is a crucial and meaningful insight for Qatra — it directly supports the premise that over-irrigation is damaging agricultural productivity.

- **Statistical Independence:** All other feature pairs show near-zero correlations (between -0.02 and 0.02), indicating that each feature contributes unique and independent information to the model.
- **No Multicollinearity Detected:** No two predictor features are highly correlated with each other, confirming that every feature contributes meaningfully to the model.
- **Implication for Clustering:** The strong negative relationship between Yield and Water_Usage is exactly the kind of pattern our unsupervised algorithms will exploit to separate efficient from inefficient farms.

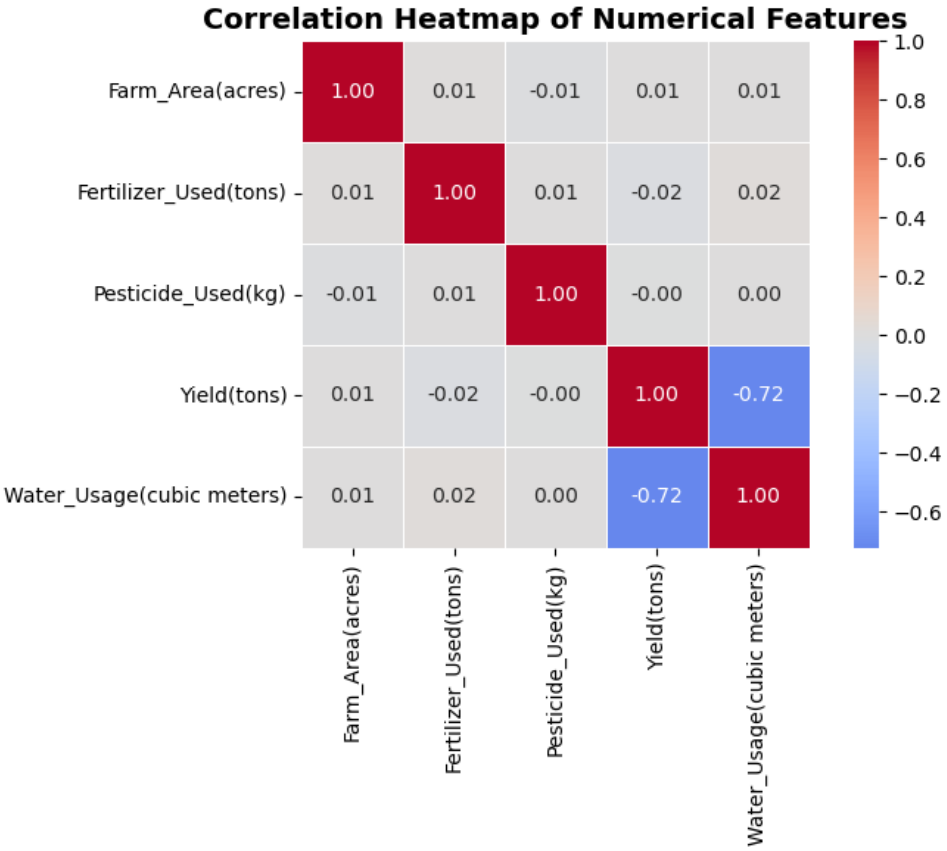

4. Categorical Data Transformation (Label Encoding)

Machine learning algorithms (including the unsupervised clustering and supervised classifiers we apply later) require all inputs to be numerical. Our dataset contains four categorical columns that must be transformed before any mathematical operations can be performed.

4.1 Identification of Categorical Features

Our dataset contains four critical non-numerical features:

- Crop_Type (10 unique values)
- Irrigation_Type (5 unique values)
- Soil_Type (5 unique values)
- Season (3 unique values)

4.2 Methodology: Label Encoding

We opted for Label Encoding which assigns a unique integer to each category. This approach was chosen over One-Hot Encoding for the following reasons:

- **Dimensionality Control:** One-Hot Encoding would expand our 9 features to 28 columns (10 + 5 + 5 + 3 dummy columns), creating high dimensionality relative to our dataset. Label Encoding keeps the feature count at 9 — no dimensionality explosion.
- **Compatibility with Tree-Based Models:** Random Forest, the primary classifier in our pipeline, splits on numerical thresholds and is not sensitive to the ordinal implications of Label Encoding.
- **Clustering Suitability:** Keeping 9 features ensures the PCA and clustering algorithms operate on a compact and meaningful feature space.

4.3 Technical Implementation

```
from sklearn.preprocessing import LabelEncoder

# Initialize the LabelEncoder
le = LabelEncoder()

# Define the categorical columns to encode
categorical_cols = ['Crop_Type', 'Irrigation_Type', 'Soil_Type', 'Season']

# Apply Label Encoding to each categorical column
# We work on a copy to preserve the original dataset for reference
dataset_encoded = dataset.copy()

for col in categorical_cols:
    dataset_encoded[col] = le.fit_transform(dataset_encoded[col])

# Display the mapping result - first 10 rows
# Farm_ID is not needed for modeling so we exclude it
print("Dataset after Label Encoding (first 10 rows):")
dataset_encoded.drop(columns=['Farm_ID']).head(10)
```

After encoding, all four categorical columns now contain integer values while the total number of columns remains unchanged at 9 features (excluding Farm_ID). The numerical columns are completely unaffected.

Dataset after Label Encoding (first 10 rows):

	Crop_Type	Farm_Area(acres)	Irrigation_Type	Fertilizer_Used(tons)	Pesticide_Used(kg)	Yield(tons)	Soil_Type	Season	Water_Usage(cubic meters)
0	4	421.36	0	8.24	1.84	44.73	2	0	30106.65
1	5	353.41	3	6.30	2.49	34.39	4	2	20471.93
2	4	260.89	1	3.92	4.39	26.65	4	0	28844.32
3	5	368.83	2	1.29	3.98	34.87	4	1	26743.29
4	4	173.73	4	2.36	2.87	38.25	4	2	47760.90
5	1	83.44	0	6.99	1.70	48.90	4	1	17255.42
6	1	296.71	0	9.31	0.79	32.67	1	1	10945.92
7	8	372.93	2	6.42	4.51	46.75	3	1	36667.05
8	6	162.13	4	9.88	3.84	47.09	3	0	19862.61
9	0	60.84	4	4.47	3.36	38.50	2	0	41171.53

5. Feature Scaling

5.1 The Need for Scaling

With all features now numerical, we must address the significant disparity in their scales before applying PCA and clustering.

- **Scale Imbalance:** Water_Usage reaches up to 95,000 cubic meters while Fertilizer_Used stays under 10 tons. Without scaling, features with larger magnitudes would completely dominate distance calculations in clustering algorithms and variance maximization in PCA.
- **PCA Prerequisite:** PCA seeks to maximize variance and requires all features to be centered on a common scale to identify the true principal components accurately.
- **Clustering Requirement:** KMeans and Agglomerative Clustering use Euclidean distance calculations, which are highly sensitive to feature scales.

5.2 Methodology: StandardScaler

We implemented the StandardScaler technique (Z-score normalization), which transforms each feature to have a mean of 0 and a standard deviation of 1. StandardScaler was chosen over MinMaxScaler because it is more robust and is the standard requirement for PCA pipelines.

5.3 Technical Implementation

```
from sklearn.preprocessing import StandardScaler
```

```
# 5.1 Feature and Target Slicing

# Before scaling we separate our feature matrix X from the Farm_ID column
# which is just an identifier and should not be used as a feature.
X = dataset_encoded.drop(columns=['Farm_ID']).values

print("Feature matrix shape:", X.shape)
print("First row before scaling:", X[0])
```

```
Feature matrix shape: (5000, 9)
First row before scaling: [4.000000e+00 4.213600e+02 0.000000e+00 8.240000e+00 1.840000e+00
 4.473000e+01 2.000000e+00 0.000000e+00 3.010665e+04]
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

After scaling, all features are centered on a common scale. The scaler object is saved for consistent application to new data during supervised classification in Phase 6.

6. Dimensionality Reduction (PCA)

The final step of Phase 2 applies Principal Component Analysis (PCA) to the scaled feature matrix. It is critical to note that PCA is applied here exclusively as an intermediate step to prepare the data for unsupervised clustering. After the target variable is generated, the supervised classifiers in Phase 6 will use X_scaled (the pre-PCA data) as input — not the PCA-reduced data.

6.1 Objectives of PCA

- **Noise Reduction:** PCA identifies the directions of maximum variance and discards directions that carry little information, helping clustering algorithms find more meaningful groupings.
- **Handling Correlation:** The strong -0.72 correlation between Yield and Water_Usage means there is redundant information. PCA will naturally compress this into a single principal component.

6.2 Component Selection

Rather than blindly choosing a fixed number of components, we first analyzed the explained variance ratio for all 9 components to make a data-driven decision.

```python
from sklearn.decomposition import PCA
import numpy as np
```

```python
# 6.1 Explained Variance Analysis

# We first fit PCA with all 9 components to see how much variance
# each principal component captures

pca_full = PCA(n_components=9)
pca_full.fit(X_scaled)

# Compute cumulative explained variance
explained_variance = pca_full.explained_variance_ratio_
cumulative_variance = np.cumsum(explained_variance)

print("Explained variance ratio per component:")
for i, (ev, cv) in enumerate(zip(explained_variance, cumulative_variance)):
    print(f"  PC{i+1}: {ev:.4f} ({ev*100:.2f}%) | Cumulative: {cv*100:.2f}%")
```

```
Explained variance ratio per component:
  PC1: 0.1914 (19.14%) | Cumulative: 19.14%
  PC2: 0.1175 (11.75%) | Cumulative: 30.89%
  PC3: 0.1151 (11.51%) | Cumulative: 42.40%
  PC4: 0.1128 (11.28%) | Cumulative: 53.68%
  PC5: 0.1118 (11.18%) | Cumulative: 64.86%
  PC6: 0.1090 (10.90%) | Cumulative: 75.77%
  PC7: 0.1069 (10.69%) | Cumulative: 86.46%
  PC8: 0.1047 (10.47%) | Cumulative: 96.93%
  PC9: 0.0307 (3.07%) | Cumulative: 100.00%
```

The variance is evenly distributed across all components — a direct consequence of the largely independent features confirmed by the correlation heatmap. Applying the standard 75% variance retention threshold, we select n_components=6, which retains 75.77% of the total variance.

6.3 Technical Implementation

```python
pca = PCA(n_components=6)
X_pca = pca.fit_transform(X_scaled)

# Verify the shape
print("Shape before PCA:", X_scaled.shape)
print("Shape after PCA:", X_pca.shape)
print(f"\nTotal variance retained: {pca.explained_variance_ratio_.sum()*100:.2f}%")
print("First row after PCA:", X_pca[0])
```

```
Shape before PCA: (5000, 9)
Shape after PCA: (5000, 6)

Total variance retained: 75.77%
First row after PCA: [ 1.74005167  1.86026357  1.10086286 -0.58685274  0.36320022 -0.254644  ]
```

7. Clustering for Target Variable Generation

The raw dataset does not contain a predefined irrigation efficiency label. To build a supervised classifier, we first engineer a target variable through unsupervised clustering. We test three algorithms — KMeans, DBSCAN, and Agglomerative Hierarchical Clustering — and evaluate each using the Silhouette Score and Davies-Bouldin Index. All clustering is performed on X_pca (the PCA-reduced data).

7.1 KMeans Clustering

Before applying KMeans we determine the optimal number of clusters k using both the Elbow Method and Silhouette Score analysis, testing values of k from 2 to 10.

- **Elbow Method Result:** The inertia decreases continuously with no sharp elbow, consistent with our uniformly distributed data.
- **Silhouette Score Result:** k=2 achieved the highest Silhouette Score of 0.2053, significantly higher than all other values which drop and flatten around 0.12-0.15.
- **Selection:** k=2 is selected as the optimal number of clusters, perfectly aligned with our binary classification goal (Efficient vs Inefficient irrigation).

```
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, davies_bouldin_score
```

```
# 7.1 KMeans Clustering

# Before applying KMeans we need to determine the optimal number of clusters k.
# We use the Elbow Method which plots the inertia (sum of squared distances
# from each point to its cluster center) for different values of k.
# The "elbow" point where inertia stops decreasing rapidly is our optimal k.

inertia = []
silhouette_scores = []
k_range = range(2, 11)  # Test k from 2 to 10

for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_pca)
    inertia.append(kmeans.inertia_)
    silhouette_scores.append(silhouette_score(X_pca, kmeans.labels_))
    print(f"k={k} | Inertia: {kmeans.inertia_:.2f} | Silhouette: {silhouette_score(X_pca, kmeans.labels_):.4f}")
```

```
k=2 | Inertia: 26814.52 | Silhouette: 0.2053
k=3 | Inertia: 24764.92 | Silhouette: 0.1469
k=4 | Inertia: 23029.19 | Silhouette: 0.1295
k=5 | Inertia: 21718.05 | Silhouette: 0.1226
k=6 | Inertia: 20596.40 | Silhouette: 0.1245
k=7 | Inertia: 19546.64 | Silhouette: 0.1262
k=8 | Inertia: 18696.83 | Silhouette: 0.1278
k=9 | Inertia: 17905.17 | Silhouette: 0.1295
k=10 | Inertia: 17242.63 | Silhouette: 0.1295
```

```python
kmeans_final = KMeans(n_clusters=2, random_state=42, n_init=10)
kmeans_labels = kmeans_final.fit_predict(X_pca)

# Evaluate the final model
kmeans_silhouette = silhouette_score(X_pca, kmeans_labels)
kmeans_db = davies_bouldin_score(X_pca, kmeans_labels)

print("KMeans Final Model (k=2) Results:")
print(f"  Silhouette Score:      {kmeans_silhouette:.4f}  (higher is better, max=1)")
print(f"  Davies-Bouldin Index:  {kmeans_db:.4f}  (lower is better, min=0)")
print(f"\nCluster distribution:")
print(f"  Cluster 0: {(kmeans_labels==0).sum()} records ({(kmeans_labels==0).sum()/len(kmeans_labels)*100:.1f}%)")
print(f"  Cluster 1: {(kmeans_labels==1).sum()} records ({(kmeans_labels==1).sum()/len(kmeans_labels)*100:.1f}%)")
```

```
KMeans Final Model (k=2) Results:
  Silhouette Score:      0.2053  (higher is better, max=1)
  Davies-Bouldin Index:  1.8659  (lower is better, min=0)

Cluster distribution:
  Cluster 0: 2382 records (47.6%)
  Cluster 1: 2618 records (52.4%)
```

7.2 DBSCAN Clustering

DBSCAN (Density-Based Spatial Clustering) was tested across a range of epsilon values (0.5 to 3.0) with min_samples=5. The algorithm fundamentally struggled with this dataset due to the uniform distribution of the agricultural data, which lacks the density variations DBSCAN requires to function effectively.

- **eps=0.5:** 99.9% of points marked as noise — epsilon too small.
- **eps=1.0:** 6 clusters found but with a negative Silhouette Score of -0.1763, indicating poor clustering quality worse than random assignment.
- **eps=1.5 and above:** Everything collapsed into a single cluster — epsilon too large.

```python
# We test a range of eps values to find the best configuration

from sklearn.cluster import DBSCAN

eps_values = [0.5, 1.0, 1.5, 2.0, 2.5, 3.0]
min_samples = 5

print("DBSCAN Parameter Search:")
print(f"{'eps':<8} {'Clusters':<12} {'Noise Points':<15} {'Noise %':<10} {'Silhouette'}")
print("-" * 60)

for eps in eps_values:
    dbscan = DBSCAN(eps=eps, min_samples=min_samples)
    dbscan_labels = dbscan.fit_predict(X_pca)

    n_clusters = len(set(dbscan_labels)) - (1 if -1 in dbscan_labels else 0)
    n_noise = (dbscan_labels == -1).sum()
    noise_pct = n_noise / len(dbscan_labels) * 100

    # Only compute silhouette if we have at least 2 clusters
    if n_clusters >= 2:
        sil = silhouette_score(X_pca, dbscan_labels)
        print(f"{eps:<8} {n_clusters:<12} {n_noise:<15} {noise_pct:<10.1f} {sil:.4f}")
    else:
        print(f"{eps:<8} {n_clusters:<12} {n_noise:<15} {noise_pct:<10.1f} {'N/A'}")
```

```
DBSCAN Parameter Search:
eps      Clusters     Noise Points    Noise %    Silhouette
------------------------------------------------------------
0.5      1            4995            99.9       N/A
1.0      6            523             10.5       -0.1763
1.5      1            0               0.0        N/A
2.0      1            0               0.0        N/A
2.5      1            0               0.0        N/A
3.0      1            0               0.0        N/A
```

```python
dbscan_best = DBSCAN(eps=1.0, min_samples=5)
dbscan_labels = dbscan_best.fit_predict(X_pca)

dbscan_silhouette = silhouette_score(X_pca, dbscan_labels)
dbscan_db = davies_bouldin_score(X_pca, dbscan_labels)

print("DBSCAN Best Result (eps=1.0, min_samples=5):")
print(f"  Silhouette Score:      {dbscan_silhouette:.4f}")
print(f"  Davies-Bouldin Index:  {dbscan_db:.4f}")
print(f"  Number of clusters:    {len(set(dbscan_labels)) - (1 if -1 in dbscan_labels else 0)}")
print(f"  Noise points:          {(dbscan_labels==-1).sum()}")
```

```
DBSCAN Best Result (eps=1.0, min_samples=5):
  Silhouette Score:      -0.1763
  Davies-Bouldin Index:  2.3898
  Number of clusters:    6
  Noise points:          523
```

7.3 Agglomerative Hierarchical Clustering

Agglomerative Clustering was tested with n_clusters=2 across three linkage methods: ward, complete, and average.

```
from sklearn.cluster import AgglomerativeClustering

linkage_methods = ['ward', 'complete', 'average']

print("Agglomerative Clustering Results (n_clusters=2):")
print(f"{'Linkage':<12} {'Silhouette':<15} {'Davies-Bouldin':<15} {'Cluster 0':<12} {'Cluster 1'}")
print("-" * 65)

agg_results = {}

for linkage in linkage_methods:
    agg = AgglomerativeClustering(n_clusters=2, linkage=linkage)
    agg_labels = agg.fit_predict(X_pca)

    sil = silhouette_score(X_pca, agg_labels)
    db = davies_bouldin_score(X_pca, agg_labels)
    c0 = (agg_labels == 0).sum()
    c1 = (agg_labels == 1).sum()

    agg_results[linkage] = {'silhouette': sil, 'db': db, 'labels': agg_labels}
    print(f"{linkage:<12} {sil:<15.4f} {db:<15.4f} {c0:<12} {c1}")
```

```
Agglomerative Clustering Results (n_clusters=2):
Linkage      Silhouette      Davies-Bouldin  Cluster 0    Cluster 1
-----------------------------------------------------------------
ward         0.2027          1.8794          2437         2563
complete     0.1903          1.9519          2692         2308
average      0.2020          1.8831          2463         2537
```

All three linkage methods produced well-balanced clusters but consistently scored slightly below KMeans on both evaluation metrics.

7.4 Algorithm Comparison & Selection

Algorithm	Silhouette Score	Davies-Bouldin Index
KMeans (k=2)	0.2053	1.8659
DBSCAN (eps=1.0)	-0.1763	2.3898
Agglomerative (ward)	0.2027	1.8794
Agglomerative (complete)	0.1903	1.9519
Agglomerative (average)	0.2020	1.8831

KMeans (k=2) is selected as the best algorithm, achieving the highest Silhouette Score (0.2053) and the lowest Davies-Bouldin Index (1.8659) across all tested configurations. DBSCAN is completely eliminated due to its negative Silhouette Score.

7.5 Target Variable Engineering (Irrigation_Efficiency)

Using the KMeans cluster assignments, we engineer the target variable Irrigation_Efficiency. To determine which cluster represents Efficient vs Inefficient irrigation, we analyze the mean Water_Usage and Yield of each cluster:

Cluster	Avg Water_Usage (m^3)	Avg Yield (tons)	Label
Cluster 0	35,258	36.59	Efficient (1)
Cluster 1	80,039	9.37	Inefficient (0)

- **Efficient (1):** Farms that use less water AND produce higher yield — consistent with the -0.72 correlation observed in the heatmap.
- **Inefficient (0):** Farms that consume significantly more water while producing lower yield — representing the over-irrigation problem Qatra aims to solve.

```
Final Irrigation_Efficiency distribution:
Irrigation_Efficiency
0    2618
1    2382
Name: count, dtype: int64

Class balance:
Irrigation_Efficiency
0    52.36
1    47.64
```

The cluster visualization (Water Usage vs Yield scatter plot) confirms perfect separation between the two groups with virtually no overlap, strongly validating the clustering approach and confirming that KMeans successfully identified meaningful and distinct groups in the data.

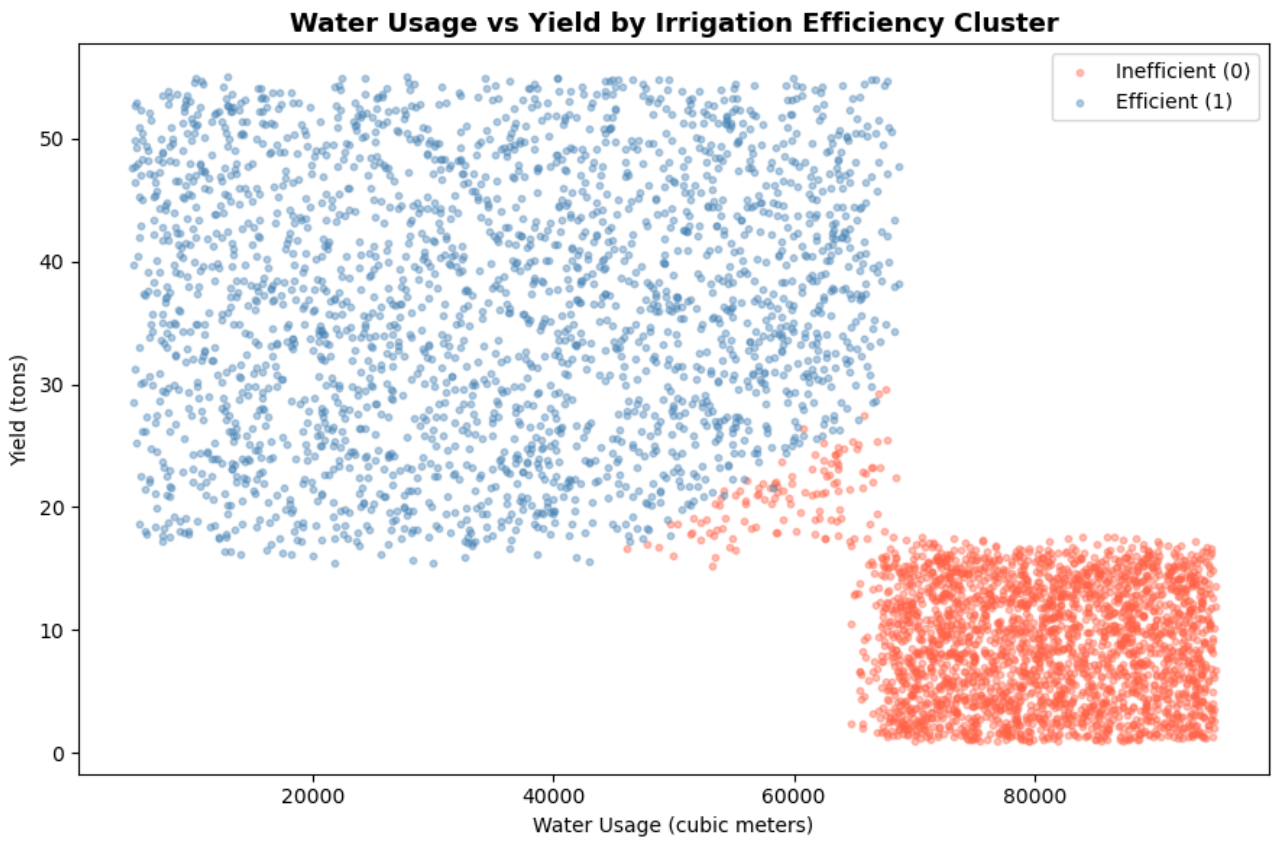

7.6 Final Data State

This marks the completion of the Pre-processing phase. The data is now fully prepared for supervised classification in Phase 6.

Step	Action	Result
EDA	Structural + Visual analysis	5,000 records, 0 nulls, 0 duplicates
Encoding	LabelEncoder on 4 categorical columns	9 features, no dimensionality explosion
Scaling	StandardScaler (Z-score)	Mean=0, Std=1 across all features
PCA	6 components retained	75.77% variance retained
Clustering	KMeans vs DBSCAN vs Agglomerative	KMeans (k=2) won on both metrics
Target Variable	Irrigation_Efficiency engineered	52.36% Inefficient / 47.64% Efficient

Phase 3: Data Warehouse Design

1. Data Warehouse Overview

The Qatra Data Warehouse is designed to support historical analysis of agricultural irrigation data through a centralized and structured storage layer. It enables interactive Power BI dashboards to explore water consumption patterns, crop yield performance, and irrigation efficiency across multiple dimensions.

The warehouse is implemented using a Star Schema — the most appropriate model for this project given its simplicity, query performance, and native compatibility with Power BI's data model and time-intelligence functions.

2. Star Schema Structure

The schema consists of one central fact table surrounded by three dimension tables:

- **Fact_Irrigation** — central table holding all measures and foreign keys (5,000 rows)
- **Dim_Date** — synthetic date dimension generated by Talend from the Season field
- **Dim_Crop** — one row per crop type (10 rows)
- **Dim_Farm** — one row per farm record (5,000 rows)

Each dimension table is linked to the fact table via a surrogate key (integer _SK column). Surrogate keys are generated during the Talend ETL process and replace natural keys in the fact table to ensure referential integrity and query performance.

3. Fact Table: Fact_Irrigation

The Fact_Irrigation table is the core of the star schema. It contains one row per farm record (5,000 rows) and stores three categories of data: foreign keys linking to each dimension, raw measures extracted directly from the dataset, and derived measures computed during the ETL process.

Field name	Data type	Key	Description
Date_SK	INT	FK	Foreign key → Dim_Date
Crop_SK	INT	FK	Foreign key → Dim_Crop
Farm_SK	INT	FK	Foreign key → Dim_Farm
Yield_Amount	FLOAT	—	Total crop output in tons
Water_Usage_Volume	FLOAT	—	Water consumed in cubic meters
Fertilizer_Usage	FLOAT	—	Fertilizer applied in tons
Pesticide_Usage	FLOAT	—	Pesticide applied in kg
Water_Use_Efficiency	FLOAT	—	Derived: Yield / Water_Usage
Water_Per_Acre	FLOAT	—	Derived: Water_Usage / Farm_Area
Yield_Per_Acre	FLOAT	—	Derived: Yield / Farm_Area
Irrigation_Efficiency_Score	INT	—	ML-generated label: 0 = Inefficient, 1 = Efficient

4. Dimension Tables

4.1 Dim_Date

Dim_Date is a synthetic dimension — it does not exist in the source dataset and is generated entirely by the Talend ETL job. Since the dataset only contains a Season field (Kharif, Rabi, or Zaid), each season is mapped to a representative calendar date range, enabling Power BI time intelligence functions (YTD, QTD, month-over-month comparisons, etc.).

Field name	Data type	Key	Description
Date_SK	**INT**	**PK**	**Surrogate key — auto-generated integer**
Full_Date	DATE	—	Calendar date mapped from Season

Field name	Data type	Key	Description
Day	INT	—	Day of the month (1–31)
Month	INT	—	Month number (1–12)
Year	INT	—	Calendar year (e.g. 2022)
Quarter	INT	—	Quarter number (1–4)
Week_Number	INT	—	ISO week number (1–52)
Month_Name	VARCHAR	—	Full month name (e.g. July)
Season_Name	VARCHAR	—	Growing season: Kharif / Rabi / Zaid

4.2 Dim_Crop

Dim_Crop holds one row per unique crop type. It enriches the raw Crop_Type field from the dataset with an agronomic category grouping, enabling higher-level analysis by crop family in Power BI.

Field name	Data type	Key	Description
Crop_SK	**INT**	**PK**	**Surrogate key — auto-generated integer**
Crop_ID	VARCHAR	NK	Natural key — sequential code (e.g. C001)
Crop_Type	VARCHAR	—	Crop name from source dataset (e.g. Potato, Rice, Maize)
Crop_Category	VARCHAR	—	Agronomic grouping: Cereal / Vegetable / Legume / Fiber / Cash Crop

4.3 Dim_Farm

Dim_Farm holds one row per farm record. It consolidates the farm's physical and operational attributes — area, soil type, and irrigation method — into a single dimension, avoiding the need for separate Dim_Soil and Dim_Irrigation tables and keeping the schema lean.

Field name	Data type	Key	Description
Farm_SK	**INT**	**PK**	**Surrogate key — auto-generated integer**
Farm_ID	VARCHAR	NK	Natural key from source dataset (e.g. SYN_0001)
Farm_Area_Acres	FLOAT	—	Farm size in acres (range: 10 to 500)
Soil_Type	VARCHAR	—	Soil classification: Loamy / Peaty / Silty / Clay / Sandy
Irrigation_Type	VARCHAR	—	Irrigation method: Drip / Manual / Flood / Rain-fed / Sprinkler
Area_Category	VARCHAR	—	Derived size bucket: Small (<167 ac) / Medium (167–333 ac) / Large (>333 ac)

5. Table Relationships

All relationships in the star schema are one-to-many (1:N), from each dimension table to the fact table. This is the standard relationship pattern for star schemas and is natively supported by Power BI's data model.

Dimension table	Cardinality	Fact table	Join key
Dim_Date	1 : N	Fact_Irrigation	Dim_Date.Date_SK = Fact_Irrigation.Date_SK
Dim_Crop	1 : N	Fact_Irrigation	Dim_Crop.Crop_SK = Fact_Irrigation.Crop_SK
Dim_Farm	1 : N	Fact_Irrigation	Dim_Farm.Farm_SK = Fact_Irrigation.Farm_SK

6. Dim_Date Generation Strategy

Since the source dataset contains no date column, a synthetic Dim_Date is generated during the Talend ETL process by mapping each record's Season value to a real calendar date range. This approach enables Power BI to recognize the date dimension as a proper time axis and activate all native time-intelligence functions.

Season	Calendar period	Months covered	Representative year
Kharif	June – October	Jun, Jul, Aug, Sep, Oct	Assigned uniformly (e.g. 2021–2023)
Rabi	November – March	Nov, Dec, Jan, Feb, Mar	Assigned uniformly (e.g. 2021–2023)
Zaid	April – June	Apr, May, Jun	Assigned uniformly (e.g. 2021–2023)

Talend implementation: *The Dim_Date table is generated using a tRowGenerator component in Talend that iterates over a predefined date range (e.g. 2021–2023), filters by season window, and assigns a Full_Date and all derived date attributes to each source record matching that season. This produces a fully populated date dimension without any real dates in the source data.*

7. Derived Measures Definition

Three measures in Fact_Irrigation are derived from raw fields rather than extracted directly from the source dataset. They are computed during the Talend ETL process before loading into the warehouse.

Measure	Formula	Business interpretation
Water_Use_Efficiency	Yield / Water_Usage	Tons of crop produced per cubic meter of water. Higher = more efficient irrigation.
Water_Per_Acre	Water_Usage / Farm_Area	Water consumption normalized by farm size. Useful for comparing farms of different scales.

Measure	Formula	Business interpretation
Yield_Per_Acre	Yield / Farm_Area	Productivity per acre. Allows fair comparison between small and large farms.

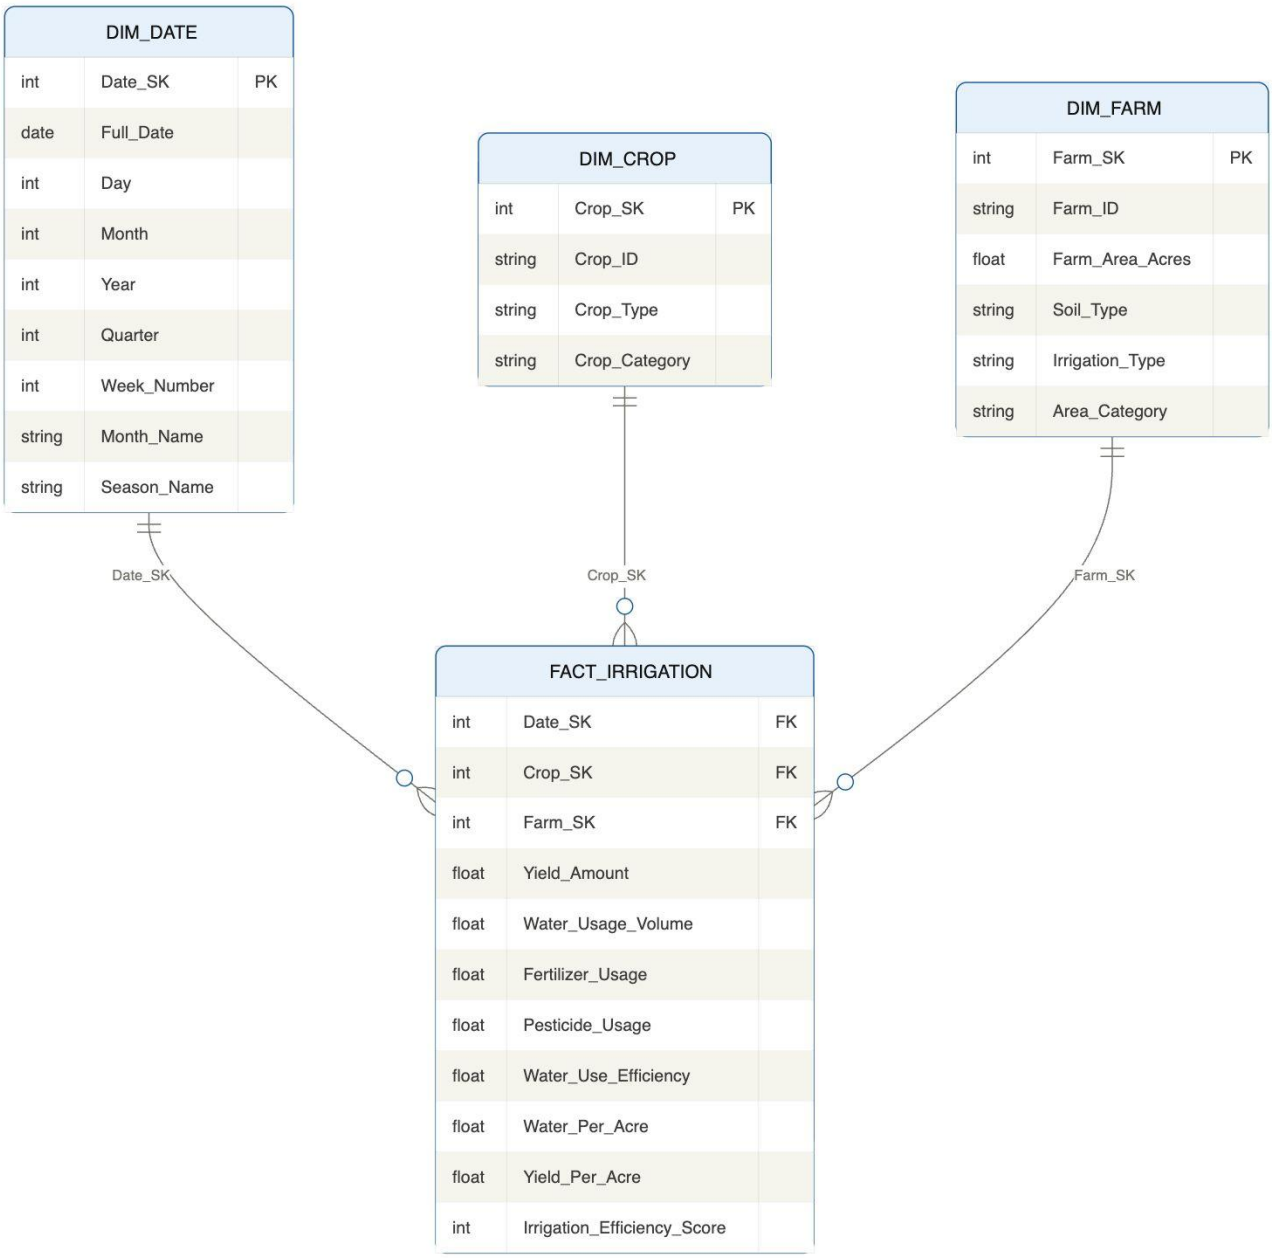

Phase 4: Feeding the data warehouse with Talend

1. Overview

This document covers Step 4 of the Qatra Decision Support System project — feeding the Data Warehouse using Talend Open Studio. The goal of this step is to extract data from the source CSV file, apply the necessary transformations, and load the results into the four tables of the Qatra star schema in PostgreSQL.

The ETL process is organized into four sequential Talend jobs, each responsible for populating one table. The execution order is strict and must be respected to ensure foreign key integrity.

Order	Job name	Target table	Source
1	**Feed_Dim_Crop**	dwh.dim_crop	Static lookup (10 rows)
2	**Feed_Dim_Farm**	dwh.dim_farm	smart_agriculture_features.csv
3	**Feed_Dim_Date**	dwh.dim_date	smart_agriculture_features.csv + generated dates
4	**Feed_Fact_Irrigation**	dwh.fact_irrigation	CSV + all three dimension tables (SK lookup)

2. Environment Setup

The following tools were used for this step:

- PostgreSQL 15 with PgAdmin 4 — serves as the relational DBMS hosting the Qatra Data Warehouse
- Talend Open Studio for Data Integration — used to design and execute the ETL jobs
- Source file: smart_agriculture_features.csv — 5,000 records, 10 columns, comma-separated

3. Database Creation (PgAdmin)

Before configuring Talend, the target database and schema were created in PostgreSQL using PgAdmin.

Steps performed:

- Opened PgAdmin and connected to the local PostgreSQL server
- Right-clicked Databases → Create → named the database DWH_QATRA
- Right-clicked DWH_QATRA → Query Tool → pasted and executed the SQL schema script

Result:

The dwh schema was created with 4 empty tables: dim_date, dim_crop, dim_farm, and fact_irrigation, each with the correct columns and constraints.

4. Talend Project & Connection Setup

Project creation

- Opened Talend Open Studio for Data Integration
- Created a new project named QATRA
- Selected the QATRA project to enter the workspace

Database connection

- In the Metadata panel, right-clicked DB Connections → Create connection
- Named the connection CONN_DWH_QATRA
- Configured the connection parameters:
    - DB Type: PostgreSQL
    - Host: localhost | Port: 5432
    - Database: DWH_QATRA | Schema: dwh
    - Username: postgres | Password: manager
- Clicked Test connection → confirmed successful connection
- Right-clicked CONN_DWH_QATRA → Retrieve schema → selected all 4 tables → Finish

5. CSV File Metadata

Before building the jobs, metadata was created in Talend to describe the source CSV file structure.

- Right-clicked File delimited in the Metadata panel → Create file delimited
- Named the metadata smart_agriculture
- Browsed to smart_agriculture_features.csv
- Set field separator to comma (,) and checked First row is header
- Verified all 10 columns were correctly detected with their types

6. Job 1 — Feed_Dim_Crop

Components used

Component	Role
tFixedFlowInput	Provides the 10 static crop rows directly inside Talend — no CSV file needed
tMap	Maps the three columns to the dim_crop output schema
tPostgresqlOutput	Writes rows to dwh.dim_crop with Insert or Update action

tFixedFlowInput data values

crop_id	crop_type	crop_category
"C001"	"Rice"	"Cereal"
"C002"	"Wheat"	"Cereal"
"C003"	"Maize"	"Cereal"
"C004"	"Barley"	"Cereal"
"C005"	"Potato"	"Vegetable"
"C006"	"Carrot"	"Vegetable"

crop_id	crop_type	crop_category
"C007"	"Tomato"	"Vegetable"
"C008"	"Soybean"	"Legume"
"C009"	"Cotton"	"Fiber"
"C010"	"Sugarcane"	"Cash Crop"

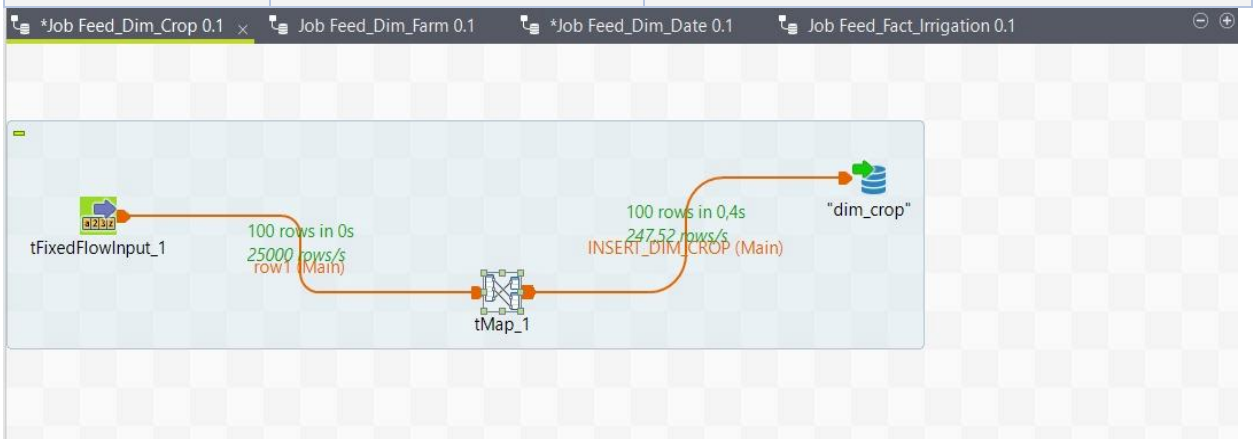

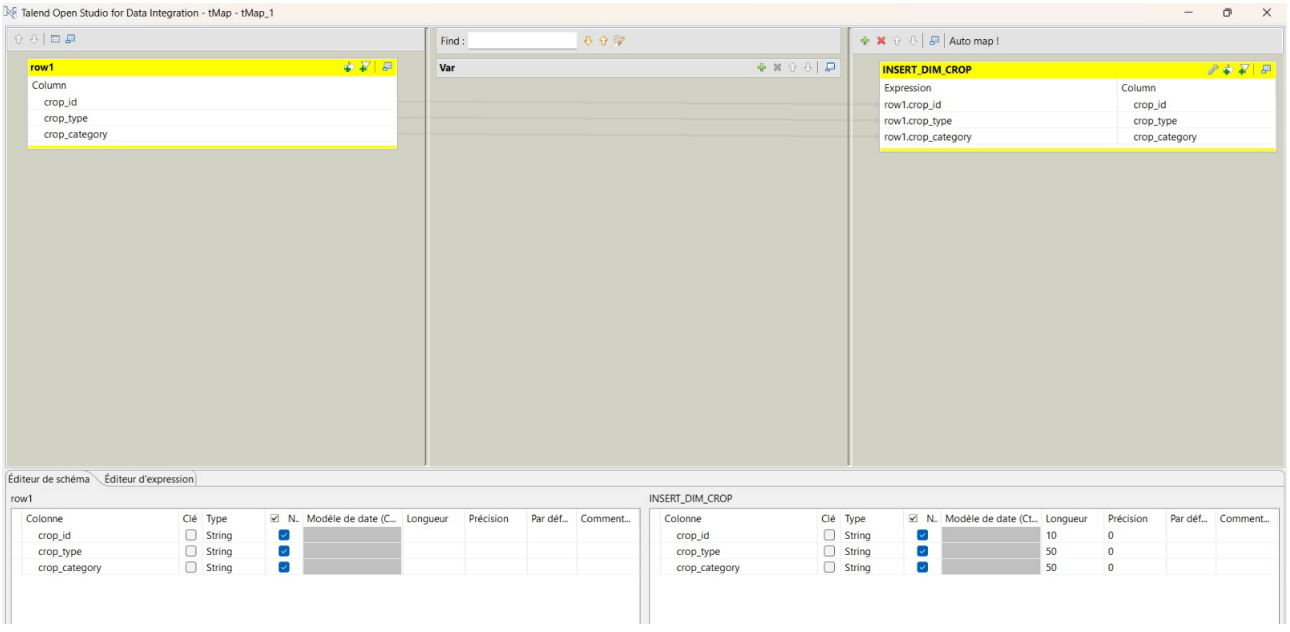

7. Job 2 — Feed_Dim_Farm

Components used

Component	Role
tFileInputDelimited	Reads all 5,000 rows from smart_agriculture_features.csv
tMap	Maps columns and derives the Area_Category field using a ternary expression
tPostgresqlOutput	Writes rows to dwh.dim_farm with Insert or Update action

Note: farm_sk is excluded from the TMap output — it is serial and auto-generated. The area_category field is derived at ETL time using a ternary expression: Small < 167 acres, Medium 167–333 acres, Large > 333 acres.

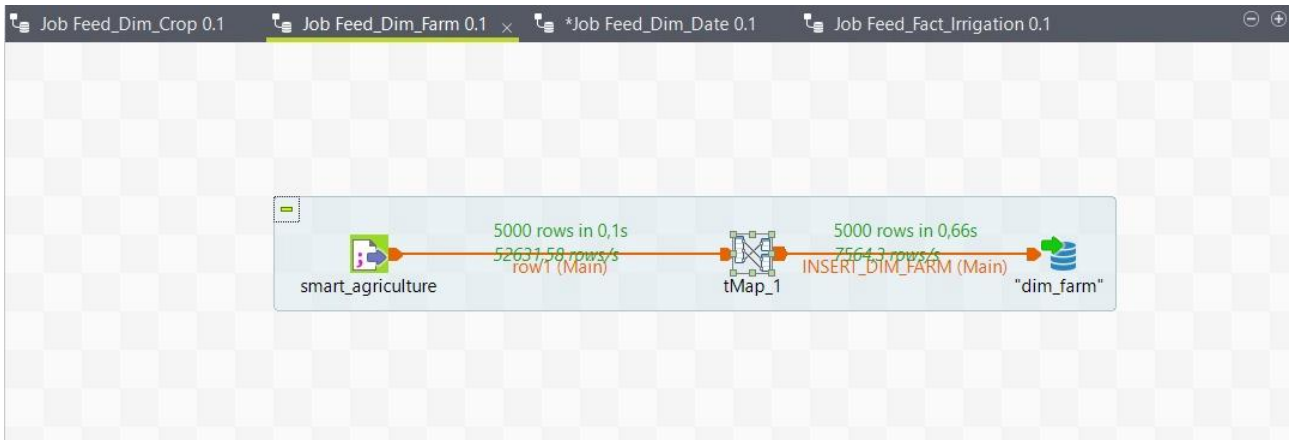

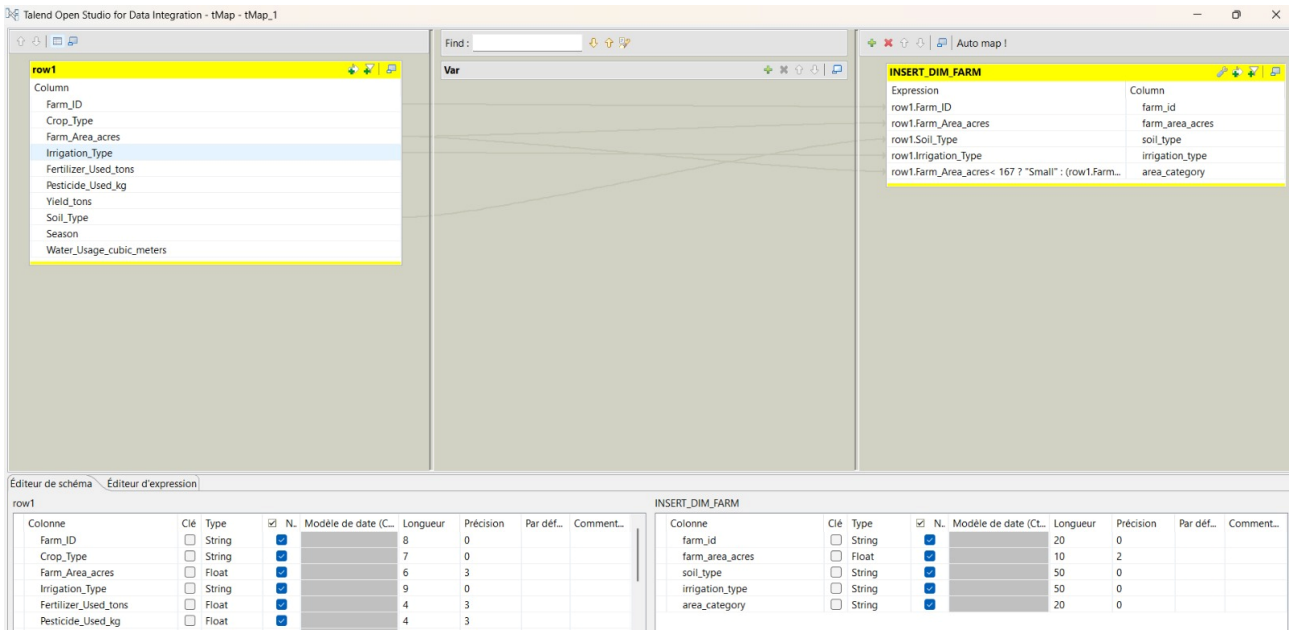

8. Job 3 — Feed_Dim_Date

The source dataset contains no date column. Dim_Date is generated synthetically by mapping each record's Season value (Kharif, Rabi, Zaid) to a representative calendar date. This enables Power BI time-intelligence functions on the resulting dashboard.

Season	Full_Date	Month	Quarter	Month_Name
Kharif	2022-07-15	7	3	July
Rabi	2022-12-15	12	4	December
Zaid	2022-05-01	5	2	May

Components used

Component	Role
tFileInputDelimited	Reads all 5,000 rows — provides the Season value for date generation
tMap	Generates all date attributes from the Season string using conditional expressions
tPostgresqlOutput	Writes rows to dwh.dim_date with Insert or Update action

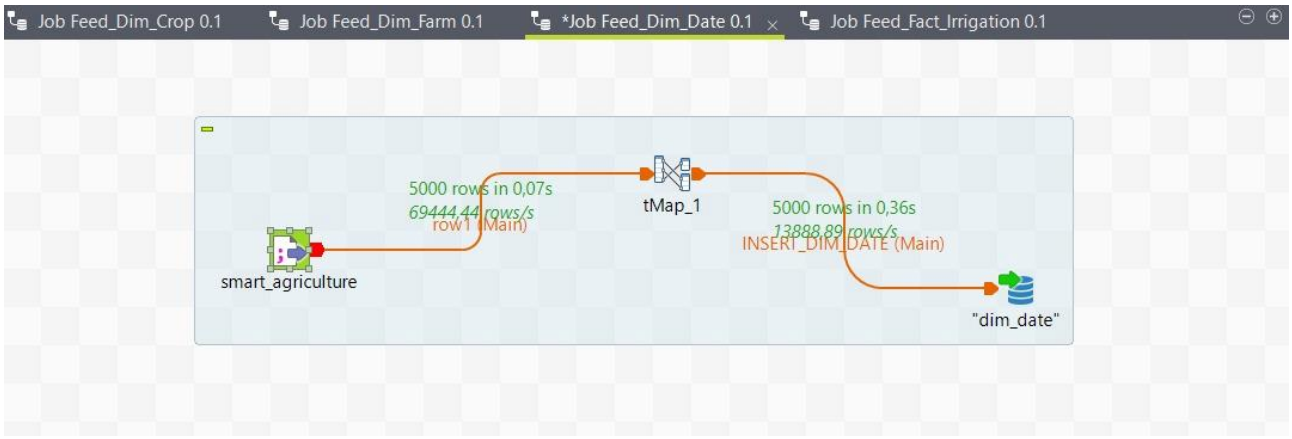

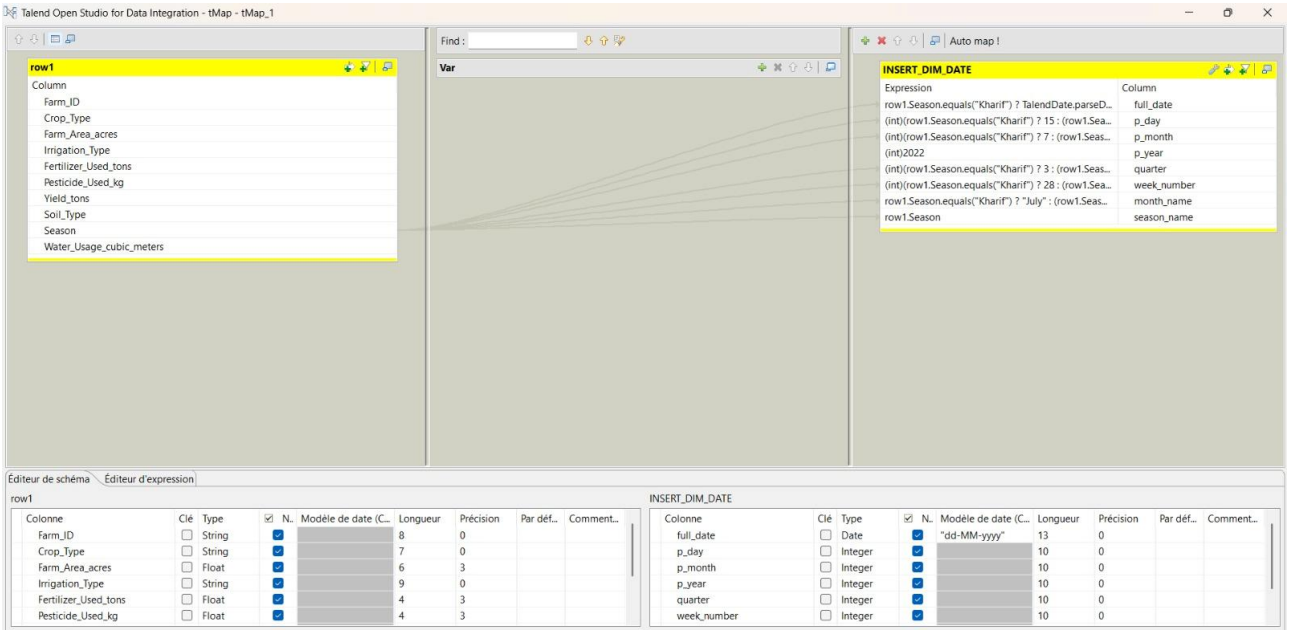

9. Job 4 — Feed_Fact_Irrigation

This is the most complex job. It reads the CSV as the main flow, looks up the surrogate keys from all three dimension tables via Inner Joins in TMap, and computes three derived measures before loading 5,000 rows into the fact table.

Components used

Component	Role
tFileInputDelimited	Main flow — reads all 5,000 rows from smart_agriculture_features.csv
tPostgresqlInput (×3)	Lookup flows — reads dim_date, dim_crop, dim_farm to resolve surrogate keys
tMap	Joins main flow with 3 lookups, computes derived measures, maps all output fields
tPostgresqlOutput	Writes rows to dwh.fact_irrigation — action: Truncate table

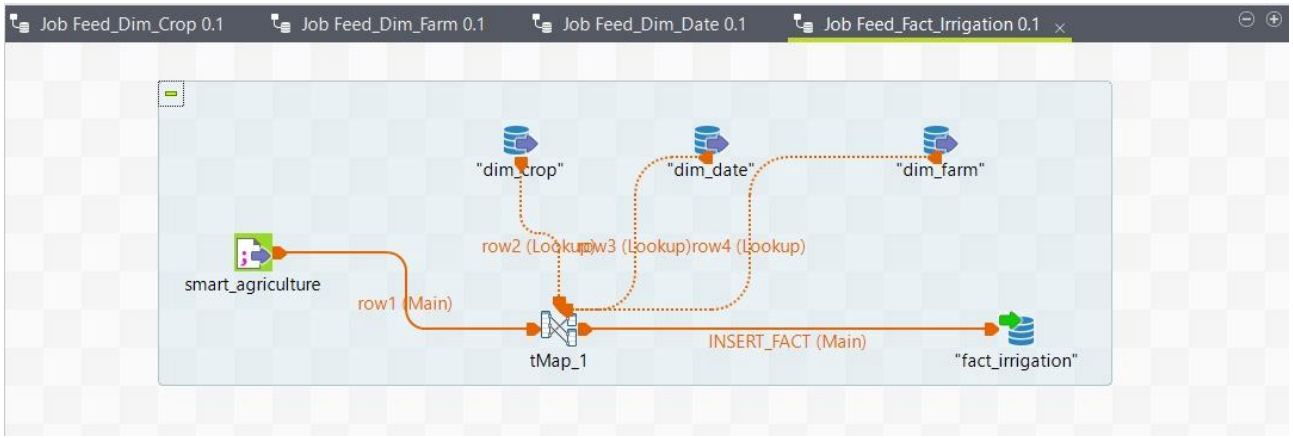
Job Feed_Dim_Crop 0.1
Job Feed_Dim_Farm 0.1
Job Feed_Dim_Date 0.1
Job Feed_Fact_Irrigation 0.1
"dim_crop"
"dim_date"
"dim_farm"
smart_agriculture
row2 (Lookup)
row3 (Lookup)
row4 (Lookup)
row1 (Main)
tMap_1
INSERT_FACT (Main)
"fact_irrigation"

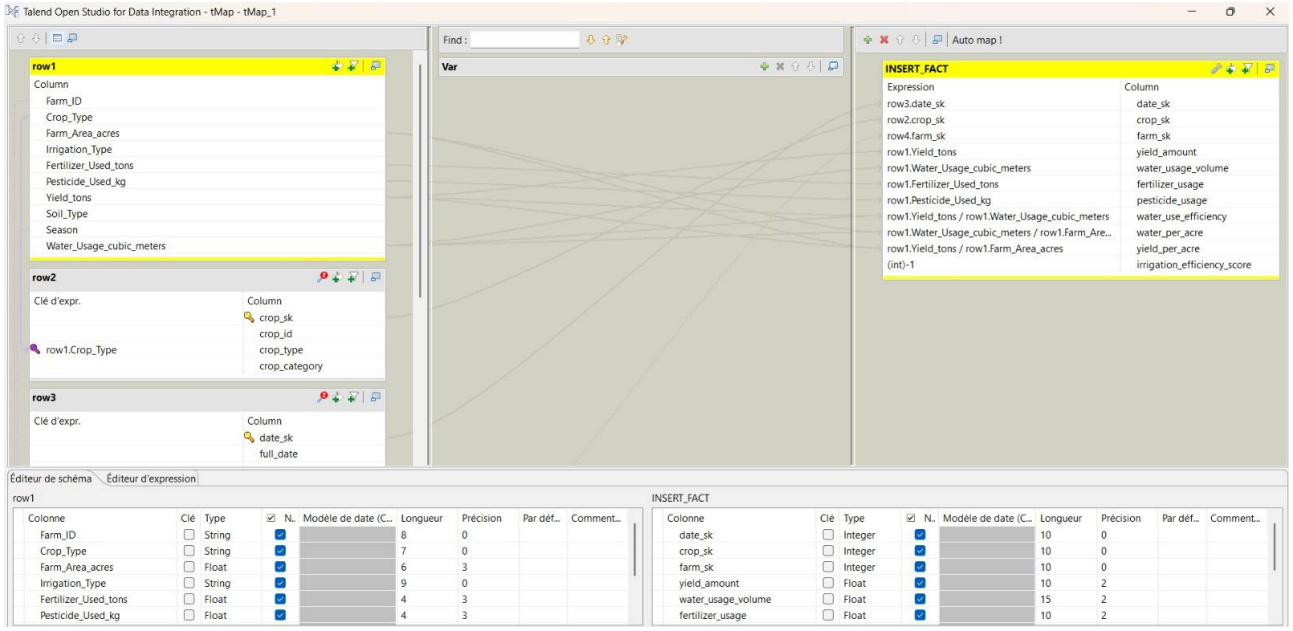
Talend Open Studio for Data Integration - tMap - tMap_1
row1
Column
Farm_ID
Crop_Type
Farm_Area_acres
Irrigation_Type
Fertilizer_Used_tons
Pesticide_Used_kg
Yield_tons
Soil_Type
Season
Water_Usage_cubic_meters
row2
Clé d'expr.
Column
crop_sk
crop_id
row1.Crop_Type
crop_type
crop_category
row3
Clé d'expr.
Column
date_sk
full_date
Find :
Var
Auto map !
INSERT_FACT
Expression
Column
row3.date_sk
date_sk
row2.crop_sk
crop_sk
row4.farm_sk
farm_sk
row1.Yield_tons
yield_amount
row1.Water_Usage_cubic_meters
water_usage_volume
row1.Fertilizer_Used_tons
fertilizer_usage
row1.Pesticide_Used_kg
pesticide_usage
row1.Yield_tons / row1.Water_Usage_cubic_meters
water_use_efficiency
row1.Water_Usage_cubic_meters / row1.Farm_Are...
water_per_acre
row1.Yield_tons / row1.Farm_Area_acres
yield_per_acre
(int)-1
irrigation_efficiency_score
Éditeur de schéma
Éditeur d'expression
row1
Colonne
Clé
Type
N..
Modèle de date (C...
Longueur
Précision
Par déf...
Comment...
Farm_ID
String
8
0
Crop_Type
String
7
0
Farm_Area_acres
Float
6
3
Irrigation_Type
String
9
0
Fertilizer_Used_tons
Float
4
3
Pesticide_Used_kg
Float
4
3
INSERT_FACT
Colonne
Clé
Type
N..
Modèle de date (C...
Longueur
Précision
Par déf...
Comment...
date_sk
Integer
10
0
crop_sk
Integer
10
0
farm_sk
Integer
10
0
yield_amount
Float
10
2
water_usage_volume
Float
15
2
fertilizer_usage
Float
10
2

Phase 5: POWER-BI

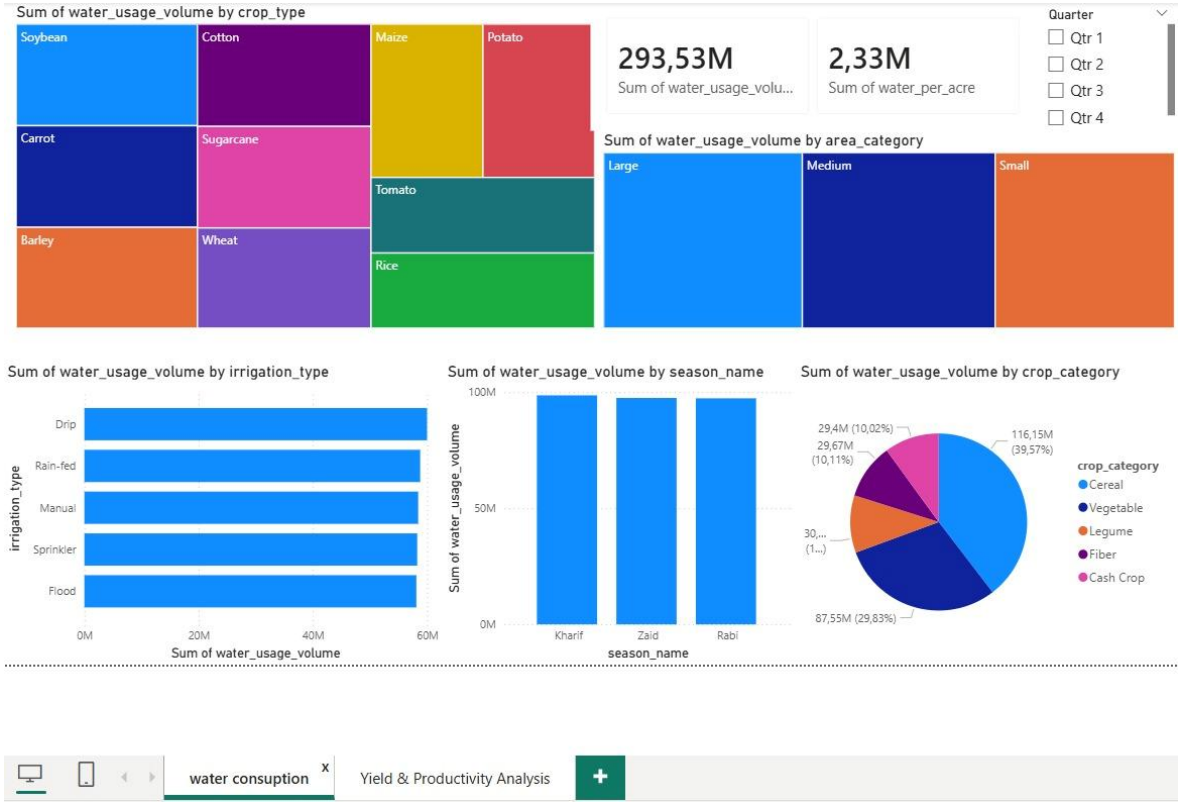

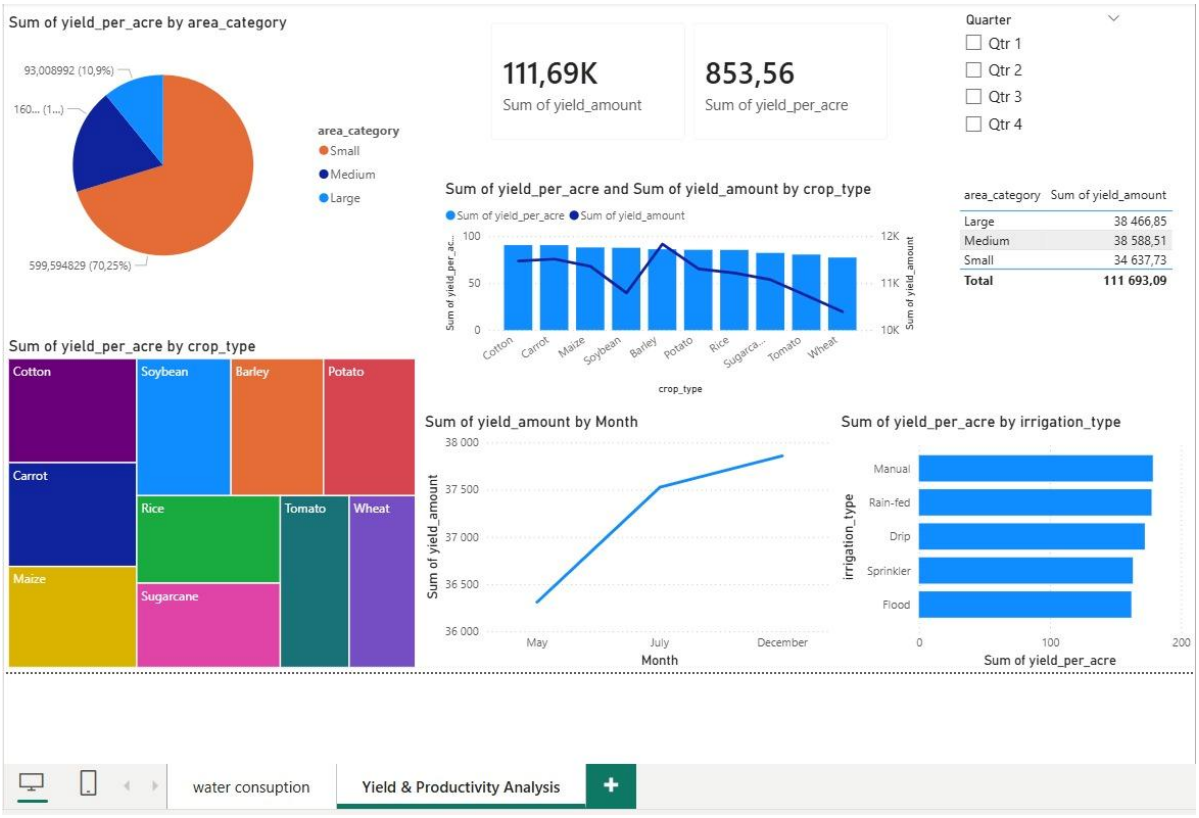

Phase 6: DATA MODELLING

In this phase we train and evaluate three supervised classification algorithms on the engineered dataset. The target variable Irrigation_Efficiency was generated in Phase 2 through KMeans clustering. All supervised models use X_scaled (the pre-PCA, 9-feature matrix) as input — not the PCA-reduced data used for clustering.

1. Dataset Partitioning (Train/Test Split)

Before training, we must reserve a portion of the data to evaluate the model's performance on unseen records.

- **Methodology:** We split the data into **80% Training** and **20% Testing**.
- **Stratification:** Crucially, we use stratify=y. This ensures that the 80/20 split maintains the same ratio in both sets. This is vital given our near-balanced but not perfectly equal class distribution.

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import (accuracy_score, precision_score,
                             recall_score, f1_score,
                             confusion_matrix, classification_report)
```

```
X = X_scaled  # Shape: (5000, 9)
y = dataset_encoded['Irrigation_Efficiency'].values  # Shape: (5000,)

print("Feature matrix shape:", X.shape)
print("Target variable shape:", y.shape)
print(f"Class distribution — Inefficient (0): {(y==0).sum()} | Efficient (1): {(y==1).sum()}")
```

```
Feature matrix shape: (5000, 9)
Target variable shape: (5000,)
Class distribution — Inefficient (0): 2618 | Efficient (1): 2382
```

```python
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)
print(f"\nTraining class distribution:")
print(f"  Inefficient (0): {(y_train==0).sum()} | Efficient (1): {(y_train==1).sum()}")
print(f"\nTesting class distribution:")
print(f"  Inefficient (0): {(y_test==0).sum()} | Efficient (1): {(y_test==1).sum()}")
```

```
Training set shape: (4000, 9)
Testing set shape: (1000, 9)

Training class distribution:
  Inefficient (0): 2094 | Efficient (1): 1906

Testing class distribution:
  Inefficient (0): 524 | Efficient (1): 476
```

2. Model Training & Evaluation

We train and evaluate three supervised classifiers, each representing a different mathematical approach to classification. Each model is evaluated using Accuracy, Precision, Recall, F1-Score, and a Confusion Matrix.

2.1 Random Forest Classifier

Random Forest is an ensemble method that combines multiple decision trees to reach a single consensus. It was selected as the primary classifier because it is highly robust against overfitting and handles non-linear relationships — which were confirmed during the correlation heatmap analysis.

- **Parameter Selection:** We empirically tested n_estimators values of 10, 50, 100, 200, and 300. Performance plateaued at n_estimators=50 with no improvement beyond this point, making it the most computationally efficient choice that achieves peak accuracy.
- **Class Weight:** class_weight='balanced' is set to account for any remaining class imbalance.

```python
#  3.1 Random Forest Classifier

# Find optimal n_estimators
from sklearn.ensemble import RandomForestClassifier

estimator_range = [10, 50, 100, 200, 300]
for n in estimator_range:
    rf_temp = RandomForestClassifier(n_estimators=n,
                                     class_weight='balanced',
                                     random_state=42)
    rf_temp.fit(X_train, y_train)
    acc = accuracy_score(y_test, rf_temp.predict(X_test))
    print(f"n_estimators={n:<6} | Accuracy: {acc:.4f}")
```

```
n_estimators=10     | Accuracy: 0.9960
n_estimators=50     | Accuracy: 0.9970
n_estimators=100    | Accuracy: 0.9970
n_estimators=200    | Accuracy: 0.9970
n_estimators=300    | Accuracy: 0.9970
```

```python
rf = RandomForestClassifier(n_estimators=50,
                            class_weight='balanced',
                            random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

print("=" * 50)
print("RANDOM FOREST RESULTS")
print("=" * 50)
print(f"Accuracy:  {accuracy_score(y_test, y_pred_rf):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_rf):.4f}")
print(f"Recall:    {recall_score(y_test, y_pred_rf):.4f}")
print(f"F1-Score:  {f1_score(y_test, y_pred_rf):.4f}")
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_rf))
print("\nClassification Report:")
print(classification_report(y_test, y_pred_rf, digits=4))
```

```
==================================================
RANDOM FOREST RESULTS
==================================================
Accuracy:  0.9970
Precision: 0.9958
Recall:    0.9979
F1-Score:  0.9969

Confusion Matrix:
[[522   2]
 [  1 475]]

Classification Report:
              precision    recall  f1-score   support

           0     0.9981    0.9962    0.9971       524
           1     0.9958    0.9979    0.9969       476

    accuracy                         0.9970      1000
   macro avg     0.9969    0.9970    0.9970      1000
weighted avg     0.9970    0.9970    0.9970      1000
```

The Random Forest achieved exceptional performance, correctly classifying 997 out of 1,000 test farms with only 3 misclassifications — 2 false positives and 1 false negative.

2.2 Support Vector Machine (SVM)

SVM finds the optimal hyperplane that maximally separates the two classes. We use the RBF (Radial Basis Function) kernel to handle the non-linear relationships confirmed during EDA.

```
from sklearn.svm import SVC

svm = SVC(kernel='rbf',
          class_weight='balanced',
          random_state=42)
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)

print("=" * 50)
print("SVM (RBF KERNEL) RESULTS")
print("=" * 50)
print(f"Accuracy:  {accuracy_score(y_test, y_pred_svm):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_svm):.4f}")
print(f"Recall:    {recall_score(y_test, y_pred_svm):.4f}")
print(f"F1-Score:  {f1_score(y_test, y_pred_svm):.4f}")
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))
print("\nClassification Report:")
print(classification_report(y_test, y_pred_svm, digits=4))
```

```
==================================================
SVM (RBF KERNEL) RESULTS
==================================================
Accuracy:  0.9900
Precision: 0.9794
Recall:    1.0000
F1-Score:  0.9896

Confusion Matrix:
[[514  10]
 [  0 476]]

Classification Report:
              precision    recall  f1-score   support

           0     1.0000    0.9809    0.9904       524
           1     0.9794    1.0000    0.9896       476

    accuracy                         0.9900      1000
   macro avg     0.9897    0.9905    0.9900      1000
weighted avg     0.9902    0.9900    0.9900      1000
```

SVM achieved a notable perfect Recall of 1.0000 — it identified every single Efficient farm in the test set with zero false negatives. However, it produced 10 false positives, slightly reducing its Precision and overall Accuracy compared to Random Forest.

2.3 K-Nearest Neighbors (KNN)

KNN classifies each point based on the majority class among its k nearest neighbors. We empirically tested k values from 3 to 15 to find the optimal configuration.

- **Parameter Selection:** k=9 achieved the highest accuracy of 0.9910 among all tested values (of k) and was selected as the optimal configuration.

```
from sklearn.neighbors import KNeighborsClassifier

# Find optimal k
k_range = [3, 5, 7, 9, 11, 13, 15]
print("KNN Parameter Search:")
for k in k_range:
    knn_temp = KNeighborsClassifier(n_neighbors=k)
    knn_temp.fit(X_train, y_train)
    acc = accuracy_score(y_test, knn_temp.predict(X_test))
    print(f"  k={k:<4} | Accuracy: {acc:.4f}")
```

```
KNN Parameter Search:
  k=3    | Accuracy: 0.9800
  k=5    | Accuracy: 0.9870
  k=7    | Accuracy: 0.9900
  k=9    | Accuracy: 0.9910
  k=11   | Accuracy: 0.9870
  k=13   | Accuracy: 0.9870
  k=15   | Accuracy: 0.9880
```

```python
knn = KNeighborsClassifier(n_neighbors=9)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)

print("=" * 50)
print("KNN RESULTS (k=9)")
print("=" * 50)
print(f"Accuracy:  {accuracy_score(y_test, y_pred_knn):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_knn):.4f}")
print(f"Recall:    {recall_score(y_test, y_pred_knn):.4f}")
print(f"F1-Score:  {f1_score(y_test, y_pred_knn):.4f}")
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_knn))
print("\nClassification Report:")
print(classification_report(y_test, y_pred_knn, digits=4))
```

```
==================================================
KNN RESULTS (k=9)
==================================================
Accuracy:  0.9910
Precision: 0.9895
Recall:    0.9916
F1-Score:  0.9906

Confusion Matrix:
[[519   5]
 [  4 472]]

Classification Report:
              precision    recall  f1-score   support

           0     0.9924    0.9905    0.9914       524
           1     0.9895    0.9916    0.9906       476

    accuracy                         0.9910      1000
   macro avg     0.9909    0.9910    0.9910      1000
weighted avg     0.9910    0.9910    0.9910      1000
```

3. Comparative Analysis & Final Model Selection

3.1 Model Performance Summary

Model	Accuracy	Precision	Recall	F1-Score
Random Forest (n=50)	0.9970	0.9958	0.9979	0.9969
SVM (RBF)	0.9900	0.9794	1.0000	0.9896
KNN (k=9)	0.9910	0.9895	0.9916	0.9906

3.2 Interpretation of Findings

- **Random Forest Superiority:** Random Forest achieved the highest performance across 3 out of 4 metrics — Accuracy (0.9970), Precision (0.9958), and F1-Score (0.9969). Its ensemble approach successfully captured the complex non-linear interactions within the agricultural data.
- **SVM Notable Recall:** SVM achieved a perfect Recall of 1.0000 — it never missed a truly Efficient farm. This is agriculturally significant as it means the system would never fail to identify an efficient farm. However, its lower Precision (0.9794) means it occasionally labels inefficient farms as efficient.
- **KNN Competitive Performance:** KNN (k=9) delivered strong and balanced results across all metrics, demonstrating that proximity-based classification is effective for this dataset despite the Label Encoding limitation noted earlier.

3.3 Final Model Selection

Based on the empirical evidence above, the Random Forest Classifier (n_estimators=50) is selected as the final engine for the Qatra Decision Support System. It provides the most stable balance between identifying Efficient farms (Recall: 0.9979) and ensuring those predictions are accurate (Precision: 0.9958), with only 3 misclassifications out of 1,000 test records.

3.4 Limitation Note

The high performance scores observed across all three classifiers are partly a consequence of the label-engineering approach. Since the target variable Irrigation_Efficiency was generated from the same features used to train the classifiers (via KMeans clustering), the supervised models are learning to reproduce a pattern that is already embedded in the data. The near-perfect cluster separation observed in the Water Usage vs Yield visualization (virtually no overlap between clusters) makes this classification task mathematically straightforward. In a real-world deployment with genuinely new and noisier farm data, performance scores would likely be lower. This is an acknowledged limitation of the unsupervised label-engineering methodology.

Conclusion

The Qatra Decision Support System successfully delivered a fully data-driven pipeline for smart irrigation management, from raw data pre-processing to actionable efficiency predictions. Through unsupervised clustering, KMeans (k=2) autonomously identified two meaningful irrigation efficiency groups, confirming that over-irrigation directly reduces crop yield — a finding strongly supported by the -0.72 correlation between Water_Usage and Yield. The supervised classification phase further validated these patterns, with the Random Forest Classifier achieving a best accuracy of 0.9970 and only 3 misclassifications out of 1,000 test records. While the high performance scores are partly a consequence of the label-engineering methodology, the system demonstrates strong potential as a practical decision support tool for farmers seeking to optimize water usage and improve agricultural productivity.